

multiple-linear-regression

November 5, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# import warnings
import warnings
warnings.filterwarnings("ignore")

# We will use some methods from the sklearn module
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split, cross_val_score
```

```
[2]: df = pd.read_csv("/kaggle/input/datasett/data (1).csv")
```

```
[3]: df.head()
```

```
[3]:
```

	Car	Model	Volume	Weight	CO2
0	Toyoty	Aygo	1000	790	99
1	Mitsubishi	Space Star	1200	1160	95
2	Skoda	Citigo	1000	929	95
3	Fiat	500	900	865	90
4	Mini	Cooper	1500	1140	105

```
[4]: df.shape
```

```
[4]: (36, 5)
```


```
[6]: print(df.describe())
```

	Volume	Weight	CO2
count	36.000000	36.000000	36.000000
mean	1611.111111	1292.277778	102.027778
std	388.975047	242.123889	7.454571
min	900.000000	790.000000	90.000000


25%	1475.000000	1117.250000	97.750000
50%	1600.000000	1329.000000	99.000000
75%	2000.000000	1418.250000	105.000000
max	2500.000000	1746.000000	120.000000

```
[7]: X = df[['Weight', 'Volume']]
     y = df['CO2']
```


```
[8]: fig, axs = plt.subplots(2, figsize = (5,5))
     plt1 = sns.boxplot(df['Weight'], ax = axs[0])
     plt2 = sns.boxplot(df['Volume'], ax = axs[1])
     plt.tight_layout()
```



```
[9]: sns.distplot(df['CO2']);
```



```
[10]: sns.pairplot(df, x_vars=['Weight', 'Volume'], y_vars='CO2', height=4, aspect=1,
      ↪ kind='scatter')
plt.show()
```



```
[12]: X_train,X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳random_state = 100)
```

```
[13]: y_train.shape
```

```
[13]: (25,)
```

```
[14]: y_test.shape
```

```
[14]: (11,)
```

```
[15]: reg_model = linear_model.LinearRegression()
```

```
[16]: reg_model = LinearRegression().fit(X_train, y_train)
```

```
[17]: print('Intercept: ',reg_model.intercept_)
# pair the feature names with the coefficients
list(zip(X, reg_model.coef_))
```

```
Intercept: 74.33882836589245
```

```
[17]: [('Weight', 0.0171800645996374), ('Volume', 0.0025046399866402976)]
```

```
[18]: y_pred= reg_model.predict(X_test)
x_pred= reg_model.predict(X_train)
```

```
[19]: print("Prediction for test set: {}".format(y_pred))
```

```
Prediction for test set: [ 90.41571939 102.16323413  99.56363213 104.56661845
101.54657652
 95.94770019 108.64011848 102.22654214  92.80374837  97.27327129
 97.57074463]
```

```
[20]: #Actual value and the predicted value
reg_model_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value':
↳y_pred})
reg_model_diff
```

```
[20]:
```

	Actual value	Predicted value
0	99	90.415719
19	105	102.163234
32	104	99.563632
35	120	104.566618
7	92	101.546577
12	99	95.947700
29	114	108.640118

33	108	102.226542
5	105	92.803748
1	95	97.273271
18	104	97.570745

```
[21]: mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = np.sqrt(metrics.mean_squared_error(y_test, y_pred))

print('Mean Absolute Error:', mae)
print('Mean Square Error:', mse)
print('Root Mean Square Error:', r2)
```

Mean Absolute Error: 6.901980901636316
Mean Square Error: 63.39765310998794
Root Mean Square Error: 7.96226432053018