

hyperparameter-tuning

December 9, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score, recall_score, \
    precision_score
from sklearn.metrics import classification_report
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A
NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy
(detected version 1.24.3
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
[2]: df = pd.read_csv('/kaggle/input/bankloan/train.csv')
df.head()
```

```
[2]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
--	----------------	---------------	-------------

0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
[3]: df = df.rename(columns=str.lower)
```

```
[4]: df.shape
```

```
[4]: (614, 13)
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_id                614 non-null    object
1   gender                 601 non-null    object
2   married                611 non-null    object
3   dependents             599 non-null    object
4   education              614 non-null    object
5   self_employed          582 non-null    object
6   applicantincome        614 non-null    int64
7   coapplicantincome      614 non-null    float64
8   loanamount             592 non-null    float64
9   loan_amount_term       600 non-null    float64
10  credit_history          564 non-null    float64
11  property_area          614 non-null    object
12  loan_status            614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
[6]: en = LabelEncoder()
cat = ['gender', 'married', 'education', '
      ↪ 'self_employed', 'property_area', 'loan_status']
for cols in cat:
    df[cols] = en.fit_transform(df[cols])
```

```
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
```

```

Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):

```

```
[7]: df['dependents'].replace('3+',3,inplace=True)
```

```
[8]: df.head()
```

```
[8]:
```

	loan_id	gender	married	dependents	education	self_employed	\
0	LP001002	1	0	0	0	0	
1	LP001003	1	1	1	0	0	
2	LP001005	1	1	0	0	1	
3	LP001006	1	1	0	1	0	
4	LP001008	1	0	0	0	0	

	applicantincome	coapplicantincome	loanamount	loan_amount_term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	credit_history	property_area	loan_status
0	1.0	2	1
1	1.0	0	0
2	1.0	2	1
3	1.0	2	1
4	1.0	2	1

```
[9]: df.isna().sum()
```

```
[9]: loan_id      0
gender          0
married         0
dependents     15
education       0
self_employed   0
applicantincome 0
coapplicantincome 0
loanamount     22
loan_amount_term 14
credit_history  50
property_area   0
loan_status     0
dtype: int64
```

```
[10]: df_clean = df
df_clean.drop('loan_id', axis=1,inplace=True)
```

```
[11]: from sklearn.impute import KNNImputer
imputer = KNNImputer(n_neighbors=3)
df_clean = pd.DataFrame(imputer.fit_transform(df),columns = df_clean.columns)
df_clean.isnull().sum()
```

/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:767:

```

FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:767:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:605:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype):
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:614:
FutureWarning: is_sparse is deprecated and will be removed in a future version.
Check `isinstance(dtype, pd.SparseDtype)` instead.
    if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):

```

```

[11]: gender          0
      married         0
      dependents      0
      education       0
      self_employed   0
      applicantincome 0
      coapplicantincome 0
      loanamount      0
      loan_amount_term 0
      credit_history   0
      property_area    0
      loan_status      0
      dtype: int64

```

```

[12]: X = df_clean.drop(columns=['loan_status']).values
      y = df_clean['loan_status'].values

```

```

[13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
      ↪random_state = 0)

```

```

[14]: sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)

```

```
[15]: rfc = RandomForestClassifier(criterion = 'entropy', random_state = 42)
      rfc.fit(X_train, y_train)
```

```
[15]: RandomForestClassifier(criterion='entropy', random_state=42)
```

```
[16]: y_pred = rfc.predict(X_test)
```

```
[17]: from sklearn import metrics
      metrix = metrics.accuracy_score(y_test, y_pred)
      print('Accuracy', metrix)
```

Accuracy 0.8051948051948052

```
[18]: from pprint import pprint
      from sklearn.model_selection import GridSearchCV, train_test_split
      from sklearn.ensemble import RandomForestClassifier
```

```
[19]: rf_classifier = RandomForestClassifier()
      print("Best Parameters:")
      pprint(rf_classifier.get_params())
```

Best Parameters:

```
{'bootstrap': True,
  'ccp_alpha': 0.0,
  'class_weight': None,
  'criterion': 'gini',
  'max_depth': None,
  'max_features': 'sqrt',
  'max_leaf_nodes': None,
  'max_samples': None,
  'min_impurity_decrease': 0.0,
  'min_samples_leaf': 1,
  'min_samples_split': 2,
  'min_weight_fraction_leaf': 0.0,
  'n_estimators': 100,
  'n_jobs': None,
  'oob_score': False,
  'random_state': None,
  'verbose': 0,
  'warm_start': False}
```

```
[20]: param_grid = {
      'n_estimators': [50, 100, 150, 200],
      'max_depth': [None, 10, 20, 30],
      'min_samples_split': [2, 5, 10],
      'min_samples_leaf': [1, 2, 4],
      'max_features': ['sqrt', 'log2'],
```

```
'bootstrap': [True, False],  
'criterion': ['gini', 'entropy']  
}
```

```
[21]: grid_search = GridSearchCV(estimator=rf_classifier, param_grid=param_grid,  
    ↪cv=5, scoring='accuracy')  
grid_search.fit(X_train, y_train)  
print("Grid Search Completed.")
```

Grid Search Completed.

```
[22]: # Train the model with the best parameters  
best_rf_model = grid_search.best_estimator_  
best_rf_model.fit(X_train, y_train)  
  
# Make predictions on the test set  
y_pred = best_rf_model.predict(X_test)
```

```
[23]: accuracy = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy)
```

Accuracy: 0.8311688311688312