

ensemble-learning

December 10, 2023

```
[104]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, \
    BaggingClassifier, ExtraTreesClassifier
from sklearn.ensemble import BaggingRegressor, RandomForestRegressor, \
    ExtraTreesRegressor
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
```

```
[105]: df = pd.read_csv("/kaggle/input/diabetes-dataset/diabetes.csv")
df.head()
```

```
[105]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[106]: df.shape
```

```
[106]: (768, 9)
```

```
[107]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

```
[108]: df.isnull().sum()
```

```
[108]: Pregnancies      0
        Glucose          0
        BloodPressure    0
        SkinThickness     0
        Insulin           0
        BMI               0
        DiabetesPedigreeFunction  0
        Age               0
        Outcome           0
        dtype: int64
```

```
[109]: df.describe()
```

```
[109]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000

max 67.100000 2.420000 81.000000 1.000000

```
[110]: categorical_val = []
        continuous_val = []
        for column in df.columns:
            if len(df[column].unique()) <= 10:
                categorical_val.append(column)
            else:
                continuous_val.append(column)
```

```
[111]: df.columns
```

```
[111]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
            'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
            dtype='object')
```

```
[112]: feature_columns = [
        'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
        'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'
    ]

    for column in feature_columns:
        print(f"column, {column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")
```

```
column, Pregnancies ==> Missing zeros : 111
column, Glucose ==> Missing zeros : 5
column, BloodPressure ==> Missing zeros : 35
column, SkinThickness ==> Missing zeros : 227
column, Insulin ==> Missing zeros : 374
column, BMI ==> Missing zeros : 11
column, DiabetesPedigreeFunction ==> Missing zeros : 0
column, Age ==> Missing zeros : 0
```

```
[113]: from sklearn.impute import SimpleImputer
        fill_values = SimpleImputer(missing_values=0, strategy="mean", copy=False)
        df[feature_columns] = fill_values.fit_transform(df[feature_columns])

        for column in feature_columns:
            print(f"column,{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")
```

```
column,Pregnancies ==> Missing zeros : 0
column,Glucose ==> Missing zeros : 0
column,BloodPressure ==> Missing zeros : 0
column,SkinThickness ==> Missing zeros : 0
column,Insulin ==> Missing zeros : 0
column,BMI ==> Missing zeros : 0
```

```
column,DiabetesPedigreeFunction ==> Missing zeros : 0
column,Age ==> Missing zeros : 0
```

```
[114]: X = df[feature_columns]
       y = df.Outcome
```

```
[115]: from sklearn.model_selection import train_test_split
       X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
       ↪random_state=42)
```

```
[116]: from sklearn.metrics import confusion_matrix, accuracy_score,
       ↪classification_report

def evaluate(model, X_train, X_test, y_train, y_test):
    y_test_pred = model.predict(X_test)
    y_train_pred = model.predict(X_train)

    print("TRAINIG RESULTS: \n=====")
    clf_report = pd.DataFrame(classification_report(y_train, y_train_pred,
    ↪output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_train, y_train_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_train, y_train_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

    print("TESTING RESULTS: \n=====")
    clf_report = pd.DataFrame(classification_report(y_test, y_test_pred,
    ↪output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_test, y_test_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_test, y_test_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")
```

Bagging Algorithms

```
[117]: from sklearn.ensemble import BaggingClassifier
       from sklearn.ensemble import RandomForestClassifier
       from sklearn.tree import DecisionTreeClassifier
       rf_clf = RandomForestClassifier(random_state=42, n_estimators=1000)
       rf_clf.fit(X_train, y_train)
       evaluate(rf_clf, X_train, X_test, y_train, y_test)
```

```
TRAINIG RESULTS:
=====
CONFUSION MATRIX:
[[349   0]
 [  0 188]]
ACCURACY SCORE:
1.0000
```

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	1.0	1.0	1.0	1.0	1.0
recall	1.0	1.0	1.0	1.0	1.0
f1-score	1.0	1.0	1.0	1.0	1.0
support	349.0	188.0	1.0	537.0	537.0

TESTING RESULTS:

=====

CONFUSION MATRIX:

```
[[123  28]
 [ 29  51]]
```

ACCURACY SCORE:

0.7532

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.809211	0.645570	0.753247	0.727390	0.752538
recall	0.814570	0.637500	0.753247	0.726035	0.753247
f1-score	0.811881	0.641509	0.753247	0.726695	0.752878
support	151.000000	80.000000	0.753247	231.000000	231.000000

```
[118]: from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

# Create a BaggingClassifier
base_classifier = DecisionTreeClassifier()
bagging_clf = BaggingClassifier(base_classifier, n_estimators=10,
    random_state=42)

# Fit the BaggingClassifier
bagging_clf.fit(X_train, y_train)

# Calculate and store accuracy scores for Bagging Classifier
bagging_scores = {
    'Train': accuracy_score(y_train, bagging_clf.predict(X_train)),
    'Test': accuracy_score(y_test, bagging_clf.predict(X_test)),
}

# Calculate and store accuracy scores for Bagging Classifier
scores = {
    'Bagging Classifier': {
        'Train': accuracy_score(y_train, bagging_clf.predict(X_train)),
        'Test': accuracy_score(y_test, bagging_clf.predict(X_test)),
    },
}
```

```
# Calculate and store accuracy scores for Random Forest
scores['Random Forest'] = {
    'Train': accuracy_score(y_train, rf_clf.predict(X_train)),
    'Test': accuracy_score(y_test, rf_clf.predict(X_test)),
}
```

Boosting Algorithms

```
[119]: from sklearn.ensemble import AdaBoostClassifier

ada_boost_clf = AdaBoostClassifier(n_estimators=30)
ada_boost_clf.fit(X_train, y_train)
evaluate(ada_boost_clf, X_train, X_test, y_train, y_test)
```

TRAINING RESULTS:

=====

CONFUSION MATRIX:

```
[[310  39]
```

```
 [ 51 137]]
```

ACCURACY SCORE:

0.8324

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.858726	0.778409	0.832402	0.818567	0.830607
recall	0.888252	0.728723	0.832402	0.808488	0.832402
f1-score	0.873239	0.752747	0.832402	0.812993	0.831056
support	349.000000	188.000000	0.832402	537.000000	537.000000

TESTING RESULTS:

=====

CONFUSION MATRIX:

```
[[123  28]
```

```
 [ 27  53]]
```

ACCURACY SCORE:

0.7619

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.820000	0.654321	0.761905	0.737160	0.762622
recall	0.814570	0.662500	0.761905	0.738535	0.761905
f1-score	0.817276	0.658385	0.761905	0.737830	0.762249
support	151.000000	80.000000	0.761905	231.000000	231.000000

```
[120]: scores['AdaBoost'] = {
    'Train': accuracy_score(y_train, ada_boost_clf.predict(X_train)),
    'Test': accuracy_score(y_test, ada_boost_clf.predict(X_test)),
}
```

```
[121]: from sklearn.ensemble import GradientBoostingClassifier

grad_boost_clf = GradientBoostingClassifier(n_estimators=100, random_state=42)
grad_boost_clf.fit(X_train, y_train)
evaluate(grad_boost_clf, X_train, X_test, y_train, y_test)
```

TRAINING RESULTS:

=====

CONFUSION MATRIX:

```
[[342   7]
```

```
 [ 19 169]]
```

ACCURACY SCORE:

0.9516

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.947368	0.960227	0.951583	0.953798	0.951870
recall	0.979943	0.898936	0.951583	0.939439	0.951583
f1-score	0.963380	0.928571	0.951583	0.945976	0.951194
support	349.000000	188.000000	0.951583	537.000000	537.000000

TESTING RESULTS:

=====

CONFUSION MATRIX:

```
[[116  35]
```

```
 [ 26  54]]
```

ACCURACY SCORE:

0.7359

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.816901	0.606742	0.735931	0.711821	0.744119
recall	0.768212	0.675000	0.735931	0.721606	0.735931
f1-score	0.791809	0.639053	0.735931	0.715431	0.738906
support	151.000000	80.000000	0.735931	231.000000	231.000000

```
[122]: scores['Gradient Boosting'] = {
        'Train': accuracy_score(y_train, grad_boost_clf.predict(X_train)),
        'Test': accuracy_score(y_test, grad_boost_clf.predict(X_test)),
    }
```

```
[123]: from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Define classifiers
log_reg = LogisticRegression(solver='liblinear')
```

```

tree = DecisionTreeClassifier()
svm_clf = SVC(gamma='scale')

estimators = [('Logistic', log_reg), ('Tree', tree), ('SVM', svm_clf)]

voting = VotingClassifier(estimators=estimators)
voting.fit(X_train, y_train)
evaluate(voting, X_train, X_test, y_train, y_test)

```

TRAINING RESULTS:

=====

CONFUSION MATRIX:

```

[[327  22]
 [ 82 106]]

```

ACCURACY SCORE:

0.8063

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.799511	0.828125	0.806331	0.813818	0.809529
recall	0.936963	0.563830	0.806331	0.750396	0.806331
f1-score	0.862797	0.670886	0.806331	0.766841	0.795610
support	349.000000	188.000000	0.806331	537.000000	537.000000

TESTING RESULTS:

=====

CONFUSION MATRIX:

```

[[131  20]
 [ 36  44]]

```

ACCURACY SCORE:

0.7576

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.784431	0.687500	0.757576	0.735966	0.750862
recall	0.867550	0.550000	0.757576	0.708775	0.757576
f1-score	0.823899	0.611111	0.757576	0.717505	0.750206
support	151.000000	80.000000	0.757576	231.000000	231.000000

```

[124]: scores['Voting'] = {
        'Train': accuracy_score(y_train, voting.predict(X_train)),
        'Test': accuracy_score(y_test, voting.predict(X_test)),
    }

```

```

[125]: scores_df = pd.DataFrame(scores)
        scores_df.plot(kind='barh', figsize=(15, 8))

```

[125]: <Axes: >

