

gaussian-naive-bayes

December 12, 2023

```
[90]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import RobustScaler
import category_encoders as ce

from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
```

```
[91]: df = pd.read_csv("/kaggle/input/adult-census-dataset/adult.csv")
df.head()
```

```
[91]:
```

	age	workclass	fnlwt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	country	salary
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[92]: df.columns
```

```
[92]: Index(['age', 'workclass', 'fnlwt', 'education', 'education-num',  
          'marital-status', 'occupation', 'relationship', 'race', 'sex',  
          'capital-gain', 'capital-loss', 'hours-per-week', 'country', 'salary'],  
          dtype='object')
```

```
[93]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32561 entries, 0 to 32560  
Data columns (total 15 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   age                   32561 non-null  int64  
1   workclass             32561 non-null  object  
2   fnlwt                 32561 non-null  int64  
3   education             32561 non-null  object  
4   education-num         32561 non-null  int64  
5   marital-status        32561 non-null  object  
6   occupation            32561 non-null  object  
7   relationship          32561 non-null  object  
8   race                  32561 non-null  object  
9   sex                   32561 non-null  object  
10  capital-gain          32561 non-null  int64  
11  capital-loss          32561 non-null  int64  
12  hours-per-week        32561 non-null  int64  
13  country               32561 non-null  object  
14  salary                32561 non-null  object  
dtypes: int64(6), object(9)  
memory usage: 3.7+ MB
```

```
[94]: df.describe()
```

```
[94]:
```

	age	fnlwt	education-num	capital-gain	capital-loss	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours-per-week
count	32561.000000
mean	40.437456

```
std      12.347429
min       1.000000
25%      40.000000
50%      40.000000
75%      45.000000
max      99.000000
```

```
[95]: #find categorical variables
categorical = [var for var in df.columns if df[var].dtype=='O']
print('There are {} categorical variables\n'.format(len(categorical)))
print('The categorical variables are :\n\n', categorical)
```

There are 9 categorical variables

The categorical variables are :

```
['workclass', 'education', 'marital-status', 'occupation', 'relationship',
'race', 'sex', 'country', 'salary']
```

```
[96]: df[categorical].head()
```

```
[96]:
```

	workclass	education	marital-status	occupation	relationship	race	sex	country	salary
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States	<=50K
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<=50K
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<=50K
3	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States	<=50K
4	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba	<=50K

```
[97]: df[categorical].columns
```

```
[97]: Index(['workclass', 'education', 'marital-status', 'occupation',
'relationship', 'race', 'sex', 'country', 'salary'],
dtype='object')
```

```
[98]: df[categorical].isnull().sum()
```

```
[98]: workclass      0
education      0
marital-status  0
occupation     0
```

```

relationship      0
race              0
sex              0
country           0
salary            0
dtype: int64

```

```
[99]: df.workclass.unique()
```

```
[99]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
            ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',
            ' Never-worked'], dtype=object)
```

```
[100]: df.workclass.value_counts()
```

```
[100]: workclass
      Private      22696
Self-emp-not-inc  2541
    Local-gov    2093
      ?         1836
    State-gov    1298
Self-emp-inc    1116
    Federal-gov   960
Without-pay      14
    Never-worked    7
Name: count, dtype: int64
```

```
[101]: df.country.unique()
```

```
[101]: array([' United-States', ' Cuba', ' Jamaica', ' India', ' ?', ' Mexico',
            ' South', ' Puerto-Rico', ' Honduras', ' England', ' Canada',
            ' Germany', ' Iran', ' Philippines', ' Italy', ' Poland',
            ' Columbia', ' Cambodia', ' Thailand', ' Ecuador', ' Laos',
            ' Taiwan', ' Haiti', ' Portugal', ' Dominican-Republic',
            ' El-Salvador', ' France', ' Guatemala', ' China', ' Japan',
            ' Yugoslavia', ' Peru', ' Outlying-US(Guam-USVI-etc)', ' Scotland',
            ' Trinidad&Tobago', ' Greece', ' Nicaragua', ' Vietnam', ' Hong',
            ' Ireland', ' Hungary', ' Holand-Netherlands'], dtype=object)
```

```
[102]: # find numerical variables
numerical = [var for var in df.columns if df[var].dtype!='O']
print('There are {} numerical variables\n'.format(len(numerical)))
print('The numerical variables are :', numerical)
```

There are 6 numerical variables

The numerical variables are : ['age', 'fnlwgt', 'education-num', 'capital-gain',

```
'capital-loss', 'hours-per-week']
```

```
[103]: df[numerical].head()
```

```
[103]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
0	39	77516	13	2174	0	40
1	50	83311	13	0	0	13
2	38	215646	9	0	0	40
3	53	234721	7	0	0	40
4	28	338409	13	0	0	40

```
[104]: df[numerical].isnull().sum()
```

```
[104]: age                0
fnlwgt                0
education-num        0
capital-gain         0
capital-loss         0
hours-per-week       0
dtype: int64
```

```
[105]: X = df.drop(['salary'], axis=1)
y = df['salary']
```

```
[106]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳ random_state = 0)
```

```
[107]: X_train.shape, X_test.shape
```

```
[107]: ((22792, 14), (9769, 14))
```

```
[108]: X_train.dtypes
```

```
[108]: age                int64
workclass            object
fnlwgt              int64
education            object
education-num        int64
marital-status       object
occupation            object
relationship         object
race                object
sex                 object
capital-gain         int64
capital-loss         int64
hours-per-week       int64
```

```
country          object
dtype: object
```

```
[109]: X_train.isnull().sum()
```

```
[109]: age          0
workclass        0
fnlwgt           0
education        0
education-num    0
marital-status   0
occupation       0
relationship     0
race             0
sex              0
capital-gain     0
capital-loss     0
hours-per-week   0
country          0
dtype: int64
```

```
[110]: X_test.isnull().sum()
```

```
[110]: age          0
workclass        0
fnlwgt           0
education        0
education-num    0
marital-status   0
occupation       0
relationship     0
race             0
sex              0
capital-gain     0
capital-loss     0
hours-per-week   0
country          0
dtype: int64
```

```
[111]: categorical
```

```
[111]: ['workclass',
'education',
'marital-status',
'occupation',
'relationship',
'race',
```

```
'sex',
'country',
'salary']
```

```
[112]: print(X_train.columns)
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
       'marital-status', 'occupation', 'relationship', 'race', 'sex',
       'capital-gain', 'capital-loss', 'hours-per-week', 'country'],
      dtype='object')
```

```
[113]: import category_encoders as ce
```

```
encoder = ce.OneHotEncoder(['workclass', 'education', 'marital_status',
                             ↪ 'occupation', 'relationship',
                             'race', 'sex', 'country'])
```

```
X_train = encoder.fit_transform(X_train)
X_test = encoder.transform(X_test)
X_train.head()
```

```
[113]:
```

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	\
32098	45	1	0	0	0	0	
25206	47	0	1	0	0	0	
23491	48	1	0	0	0	0	
12367	29	1	0	0	0	0	
7054	23	1	0	0	0	0	

	workclass_6	workclass_7	workclass_8	workclass_9	...	country_33	\
32098	0	0	0	0	...	0	
25206	0	0	0	0	...	0	
23491	0	0	0	0	...	0	
12367	0	0	0	0	...	0	
7054	0	0	0	0	...	0	

	country_34	country_35	country_36	country_37	country_38	country_39	\
32098	0	0	0	0	0	0	
25206	0	0	0	0	0	0	
23491	0	0	0	0	0	0	
12367	0	0	0	0	0	0	
7054	0	0	0	0	0	0	

	country_40	country_41	country_42
32098	0	0	0
25206	0	0	0
23491	0	0	0
12367	0	0	0

7054 0 0 0

[5 rows x 108 columns]

```
[114]: cols = X_train.columns
```

```
[115]: from sklearn.preprocessing import RobustScaler
scaler = RobustScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[116]: X_train = pd.DataFrame(X_train, columns=[cols])
X_test = pd.DataFrame(X_test, columns=[cols])
X_train.head()
```

```
[116]:      age workclass_1 workclass_2 workclass_3 workclass_4 workclass_5 \
0  0.40          0.0          0.0          0.0          0.0          0.0
1  0.50         -1.0          1.0          0.0          0.0          0.0
2  0.55          0.0          0.0          0.0          0.0          0.0
3 -0.40          0.0          0.0          0.0          0.0          0.0
4 -0.70          0.0          0.0          0.0          0.0          0.0

      workclass_6 workclass_7 workclass_8 workclass_9 ... country_33 country_34 \
0          0.0          0.0          0.0          0.0 ...          0.0          0.0
1          0.0          0.0          0.0          0.0 ...          0.0          0.0
2          0.0          0.0          0.0          0.0 ...          0.0          0.0
3          0.0          0.0          0.0          0.0 ...          0.0          0.0
4          0.0          0.0          0.0          0.0 ...          0.0          0.0

      country_35 country_36 country_37 country_38 country_39 country_40 \
0          0.0          0.0          0.0          0.0          0.0          0.0
1          0.0          0.0          0.0          0.0          0.0          0.0
2          0.0          0.0          0.0          0.0          0.0          0.0
3          0.0          0.0          0.0          0.0          0.0          0.0
4          0.0          0.0          0.0          0.0          0.0          0.0

      country_41 country_42
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0
```

[5 rows x 108 columns]

Gaussian Naive Bayes classifier


```
[117]: # train a Gaussian Naive Bayes classifier on the training set
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)
```

```
[117]: GaussianNB()
```

```
[118]: y_pred = gnb.predict(X_test)
y_pred
```

```
[118]: array([' <=50K', ' <=50K', ' >50K', ..., ' >50K', ' <=50K', ' <=50K'],
      dtype='<U6')
```

```
[119]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.93	0.79	0.86	7407
>50K	0.56	0.81	0.66	2362
accuracy			0.80	9769
macro avg	0.74	0.80	0.76	9769
weighted avg	0.84	0.80	0.81	9769

```
[120]: from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

accuracy = accuracy_score(y_test, y_pred)
print('Model accuracy score: {0:0.4f}'.format(accuracy))
cm = confusion_matrix(y_test, y_pred)
print('Confusion matrix\n\n', cm)
```

Model accuracy score: 0.7973

Confusion matrix

```
[[5871 1536]
 [ 444 1918]]
```

```
[121]: sns.heatmap(cm, annot=True, fmt='d', cmap="bwr")
```

```
[121]: <Axes: >
```

